

Shuyi Yu

✉ yu.shuyi.yx3@naist.ac.jp
🌐 <https://yyushuyui.github.io>

🎓 Google Scholar
📖 Chinese



Personal Profile

I am Shuyi Yu, a Master's student at the Nara Institute of Science and Technology (NAIST), supervised by Prof. Taro Watanabe, continuing to a PhD in the same lab from October 2026. From November 2026, I will also work as a Research Assistant at the Research and Development Center for Large Language Models (LLMC), National Institute of Informatics (NII).

My work is driven by one question: why do language models behave in ways their developers did not intend, and what can be done about it. I approach this from both ends, studying how safety alignment breaks down under adversarial input, and how post-training decisions shape the behavior that alignment is supposed to constrain. My goal is to make the behavior of language models something that can be understood, predicted, and deliberately controlled rather than merely observed.

Education

- 2026.10 – **Ph.D. in Information Science**, Nara Institute of Science and Technology, Nara, Japan.
NLP Laboratory, Supervisor: Prof. Taro Watanabe.
- 2024.10 – 2026.9 **Master in Information Science**, Nara Institute of Science and Technology, Nara, Japan.
NLP Laboratory, Supervisor: Prof. Taro Watanabe.
Thesis title: *Random Character-Level Perturbations Amplify LLM Jailbreak Attacks*.
- 2019.9 – 2023.6 **Bachelor in English Language and Literature**, Hangzhou City University, Hangzhou, China.
Thesis title: *Small Business English Corpus Construction and Writing Teaching Application*

Research Interests

- Safety & Ethics **■** Safety under jailbreak attack scenarios; regulatory compliance of LLM outputs; bias and prejudice in model behavior.
- Personalized NLP **■** How LLM outputs vary in style, and how to model individual user preferences over such outputs.
- Post-training **■** The broader effects and value of post-training methods (SFT/DPO/RLHF/RLVR) on model behavior.

Publications

Journal Articles

- 1 **S. Yu** et al., “Random character-level perturbations amplify LLM jailbreak attacks,” *Transactions on Machine Learning Research*, 2026. 🔗 URL: <https://openreview.net/forum?id=BXsOIppKEI>

Under Review

- 1 **S. Yu**, Y. Sakai, H. Kamigaito, and T. Watanabe, “From lexical to entity-level: Predicting individual knowledge boundaries via knowledge graph label propagation,” Submitted to ACL-IJCNLP 2026, 2026.

Employment History

2026.11 –  **Research Assistant (incoming)**, Research and Development Center for Large Language Models (LLMC), National Institute of Informatics (NII), Remote (based in Nara, Japan), starting 2026.11.1.
Distillation bias detection and mitigation: identifying and mitigating the behavioral and knowledge biases that student models inherit from teacher models through knowledge distillation.

Skills

Programming	 Python, Bash, Pandas, NumPy.
Frameworks	 PyTorch, HuggingFace (Transformers/PEFT/Datasets/TRL), LLaMA Factory.
Training & Alignment	 SFT, LoRA/QLoRA, DPO, data preprocessing.
Inference & APIs	 vLLM, Ollama, OpenAI API.
MLOps & Tools	 Linux, Git, SLURM, Weights & Biases, $\LaTeX$.
Languages	 Chinese (native), English (business), Japanese (daily conversation).

Miscellaneous Experience

Certifications

2022.7  **JLPT N1** (135/180).
2021.9  **TOEIC** (960/990).